

```
In [72]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import scipy.stats as st
import warnings
import statsmodels.api as sm
import statsmodels.formula.api as smf
warnings.filterwarnings('ignore')
```

Loading Dataset

```
In [12]: # encoding='unicode_escape' is used for removing unwanted spaces or _ - in the csv path
df=pd.read_csv(r'C:\Users\dharm\Downloads\2020_Yellow_Taxi_Trip_Data.csv',encoding='unicode_escape')
```

```
In [19]: df.shape
```

```
Out[19]: (24648499, 18)
```

```
In [13]: df.head()
```

```
Out[13]:
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	payme
0	1.0	01/01/2020 12:28:15 AM	01/01/2020 12:33:03 AM	1.0	1.2	1.0	N	238	239	
1	1.0	01/01/2020 12:35:39 AM	01/01/2020 12:43:04 AM	1.0	1.2	1.0	N	239	238	
2	1.0	01/01/2020 12:47:41 AM	01/01/2020 12:53:52 AM	1.0	0.6	1.0	N	238	238	
3	1.0	01/01/2020 12:55:23 AM	01/01/2020 01:00:14 AM	1.0	0.8	1.0	N	238	151	
4	2.0	01/01/2020 12:01:58 AM	01/01/2020 12:04:16 AM	1.0	0.0	1.0	N	193	193	

```
In [81]: df_mean = np.mean(df['fare_amount'])
df_mean
```

```
Out[81]: 8.72986959585334
```

```
In [82]: df_std = np.std(df['fare_amount'])
df_std
```

```
Out[82]: 3.595938337341993
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24648499 entries, 0 to 24648498
Data columns (total 18 columns):
#   Column              Dtype
---  -
0   VendorID            float64
1   tpep_pickup_datetime object
2   tpep_dropoff_datetime object
3   passenger_count     float64
4   trip_distance       float64
5   RatecodeID          float64
6   store_and_fwd_flag  object
7   PULocationID        int64
8   DOLocationID        int64
9   payment_type        float64
10  fare_amount         float64
11  extra               float64
12  mta_tax             float64
13  tip_amount          float64
14  tolls_amount        float64
15  improvement_surcharge float64
16  total_amount        float64
17  congestion_surcharge float64
dtypes: float64(13), int64(2), object(3)
memory usage: 3.3+ GB
```

EDA

```
In [6]: df['tpep_pickup_datetime'] = pd.to_datetime(df['tpep_pickup_datetime'])
df['tpep_dropoff_datetime'] = pd.to_datetime(df['tpep_dropoff_datetime'])

-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-6-48671b08eba7> in <module>
----> 1 df['tpep_pickup_datetime'] = pd.to_datetime(df['tpep_pickup_datetime'])
      2 df['tpep_dropoff_datetime'] = pd.to_datetime(df['tpep_dropoff_datetime'])

~\anaconda3\lib\site-packages\pandas\core\timeseries.py in to_datetime(arg, errors, dayfirst, yearfirst, utc, format, exact, unit, infer_datetime_format, origin, cache)
    803         result = arg.map(cache_array)
    804     else:
--> 805         values = convert_listlike(arg._values, format)
    806         result = arg._constructor(values, index=arg.index, name=arg.name)
    807     elif isinstance(arg, (ABCDDataFrame, abc.MutableMapping)):

~\anaconda3\lib\site-packages\pandas\core\timeseries.py in _convert_listlike_datetimes(arg, format, name, tz, unit, errors, infer_datetime_format, dayfirst, yearfirst, exact)
    463         assert format is None or infer_datetime_format
    464         utc = tz == "utc"
--> 465         result, tz_parsed = objects_to_datetime64ns(
    466             arg,
    467             dayfirst=dayfirst,

~\anaconda3\lib\site-packages\pandas\core\arrays\timeseries.py in objects_to_datetime64ns(data, dayfirst, yearfirst, utc, errors, require_iso8601, allow_object)
    2073
    2074     try:
-> 2075         result, tz_parsed = tslib.array_to_datetime(
    2076             data,
    2077             errors=errors,

pandas\_libs\tslib.pyx in pandas._libs.tslib.array_to_datetime()

pandas\_libs\tslib.pyx in pandas._libs.tslib.array_to_datetime()

pandas\_libs\tslib\parsing.pyx in pandas._libs.tslib.parsing.parse_datetime_string()

~\anaconda3\lib\site-packages\dateutil\parser\_parser.py in parse(timestr, parserinfo, **kwargs)
    1372         return parser(parserinfo).parse(timestr, **kwargs)
    1373     else:
-> 1374         return DEFAULTPARSER.parse(timestr, **kwargs)
    1375
    1376

~\anaconda3\lib\site-packages\dateutil\parser\_parser.py in parse(self, timestr, default, ignoretz, tzinfos, **kwargs)
    644         second=0, microsecond=0)
    645
-> 646         res, skipped_tokens = self._parse(timestr, **kwargs)
    647
    648         if res is None:

~\anaconda3\lib\site-packages\dateutil\parser\_parser.py in _parse(self, timestr, dayfirst, yearfirst, fuzzy, fuzzy_with_tokens)
    870         return None, None
    871
-> 872         if not info.validate(res):
    873             return None, None
    874

~\anaconda3\lib\site-packages\dateutil\parser\_parser.py in validate(self, res)
    390
    391         if ((res.tzoffset == 0 and not res.tzname) or
-> 392             (res.tzname == 'Z' or res.tzname == 'z')):
    393             res.tzname = "UTC"
    394             res.tzoffset = 0

KeyboardInterrupt:
```

```
In [7]: df.dtypes

Out[7]: VendorID                float64
tpep_pickup_datetime          object
tpep_dropoff_datetime         object
passenger_count              float64
trip_distance                float64
RatecodeID                   float64
store_and_fwd_flag            object
PULocationID                  int64
DOLocationID                  int64
payment_type                  float64
fare_amount                   float64
extra                         float64
mta_tax                       float64
tip_amount                    float64
tolls_amount                  float64
improvement_surcharge         float64
total_amount                  float64
congestion_surcharge          float64
dtype: object

In [ ]: df['duration'] = [df['tpep_dropoff_datetime'] - df['tpep_pickup_datetime']].dt.total_seconds()/60
df['duration']

In [14]: df = df[['passenger_count', 'payment_type', 'fare_amount', 'trip_distance', 'tpep_pickup_datetime', 'tpep_dropoff_datetime']]
df

Out[14]:
```

	passenger_count	payment_type	fare_amount	trip_distance	tpep_pickup_datetime	tpep_dropoff_datetime
0	1.0	1.0	6.00	1.20	01/01/2020 12:28:15 AM	01/01/2020 12:33:03 AM
1	1.0	1.0	7.00	1.20	01/01/2020 12:35:39 AM	01/01/2020 12:43:04 AM
2	1.0	1.0	6.00	0.60	01/01/2020 12:47:41 AM	01/01/2020 12:53:52 AM
3	1.0	1.0	5.50	0.80	01/01/2020 12:55:23 AM	01/01/2020 01:00:14 AM
4	1.0	2.0	3.50	0.00	01/01/2020 12:01:58 AM	01/01/2020 12:04:16 AM
...
24648494	NaN	NaN	32.49	9.22	12/31/2020 11:44:35 PM	01/01/2021 12:01:22 AM
24648495	NaN	NaN	13.22	4.79	12/31/2020 11:41:36 PM	12/31/2020 11:50:32 PM
24648496	NaN	NaN	69.31	28.00	12/31/2020 11:01:17 PM	12/31/2020 11:40:37 PM
24648497	NaN	NaN	35.95	7.08	12/31/2020 11:31:29 PM	12/31/2020 11:44:22 PM
24648498	NaN	NaN	17.09	2.35	12/31/2020 11:12:48 PM	12/31/2020 11:24:51 PM

24648499 rows x 6 columns

```
In [15]: df['tpep_pickup_datetime'] = pd.to_datetime(df['tpep_pickup_datetime'])
df['tpep_dropoff_datetime'] = pd.to_datetime(df['tpep_dropoff_datetime'])
```

```
In [21]: df['duration'] = [df['tpep_pickup_datetime'] - df['tpep_dropoff_datetime']]
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-21-0edbe26973a8> in <module>
----> 1 df['duration'] = [df['tpep_pickup_datetime'] - df['tpep_dropoff_datetime']]

~\anaconda3\lib\site-packages\pandas\core\frame.py in __setitem__(self, key, value)
   3161     else:
   3162         # set column
-> 3163         self._set_item(key, value)
   3164
   3165     def _setitem_slice(self, key: slice, value):

~\anaconda3\lib\site-packages\pandas\core\frame.py in _set_item(self, key, value)
   3240     """
   3241     self._ensure_valid_index(value)
-> 3242     value = self._sanitize_column(key, value)
   3243     NDFrame._set_item(self, key, value)
   3244

~\anaconda3\lib\site-packages\pandas\core\frame.py in _sanitize_column(self, key, value, broadcast)
   3897
   3898     # turn me into an ndarray
-> 3899     value = sanitize_index(value, self.index)
   3900     if not isinstance(value, (np.ndarray, Index)):
   3901         if isinstance(value, list) and len(value) > 0:

~\anaconda3\lib\site-packages\pandas\core\internals\construction.py in sanitize_index(data, index)
    749     """
    750     if len(data) != len(index):
-> 751         raise ValueError(
    752             "Length of values "
    753             f"({len(data)}) "

ValueError: Length of values (1) does not match length of index (24648499)
```

```
In [22]: df.isnull().sum()
```

```
Out[22]: passenger_count    809568
         payment_type      809568
         fare_amount         0
         trip_distance       0
         tpep_pickup_datetime 0
         tpep_dropoff_datetime 0
         dtype: int64
```

```
In [23]: (809568/len(df) * 100)
```

```
Out[23]: 3.2844515197456854
```

```
In [24]: df.dropna(inplace = True) #inplace = True will save the operation in the dataframe (just like save button in the excel)
         df
```

```
Out[24]:
```

	passenger_count	payment_type	fare_amount	trip_distance	tpep_pickup_datetime	tpep_dropoff_datetime
0	1.0	1.0	6.0	1.20	2020-01-01 00:28:15	2020-01-01 00:33:03
1	1.0	1.0	7.0	1.20	2020-01-01 00:35:39	2020-01-01 00:43:04
2	1.0	1.0	6.0	0.60	2020-01-01 00:47:41	2020-01-01 00:53:52
3	1.0	1.0	5.5	0.80	2020-01-01 00:55:23	2020-01-01 01:00:14
4	1.0	2.0	3.5	0.00	2020-01-01 00:01:58	2020-01-01 00:04:16
...
24549234	1.0	2.0	33.0	11.30	2020-12-31 23:05:33	2020-12-31 23:31:36
24549235	1.0	1.0	9.0	2.18	2020-12-31 22:57:20	2020-12-31 23:05:33
24549236	1.0	1.0	9.5	2.52	2020-12-31 23:40:35	2020-12-31 23:48:43
24549237	1.0	1.0	4.5	0.59	2020-12-31 23:54:57	2020-12-31 23:57:39
24549238	1.0	2.0	18.5	6.06	2020-12-31 23:11:16	2020-12-31 23:24:08

23838931 rows × 6 columns

```
In [28]: df['passenger_count'] = df['passenger_count'].astype('int64') #astype will convert tye datatype
         df['payment_type'] = df['payment_type'].astype('int64') #astype will convert tye datatype
```

```
In [26]: df[df.duplicated()]
```

```
Out[26]:
```

	passenger_count	payment_type	fare_amount	trip_distance	tpep_pickup_datetime	tpep_dropoff_datetime
39458	1	1	13.5	3.0	2020-01-01 02:29:50	2020-01-01 02:47:07
561748	1	1	4.0	0.4	2020-01-04 11:07:40	2020-01-04 11:10:37
967243	1	1	5.0	0.8	2020-01-06 16:50:12	2020-01-06 16:54:17
1060000	1	1	6.0	1.0	2020-01-07 08:22:23	2020-01-07 08:28:24
1845592	1	1	4.0	0.5	2020-01-10 19:28:12	2020-01-10 19:31:00
...
22084090	1	1	5.5	0.7	2020-11-09 13:24:06	2020-11-09 13:30:09
23388486	1	1	5.5	1.0	2020-12-04 19:32:16	2020-12-04 19:37:11
24232098	1	1	5.5	0.7	2020-12-23 10:49:44	2020-12-23 10:55:31
24240232	1	1	11.0	2.0	2020-12-23 12:34:30	2020-12-23 12:49:18
24520705	1	1	5.0	0.8	2020-12-31 12:21:11	2020-12-31 12:25:31

66 rows × 6 columns

```
In [29]: df.drop_duplicates(inplace=True)
```

```
In [30]: df['passenger_count'].value_counts(normalize = True) #value_counts gives frequency for the categories and how much time it is pr
```

```
Out[30]: 1    0.734563
2    0.140490
3    0.036607
5    0.031533
0    0.020529
6    0.019906
4    0.016364
7    0.000004
8    0.000002
9    0.000002
Name: passenger_count, dtype: float64
```

```
In [31]: df['payment_type'].value_counts(normalize = True)
```

```
Out[31]: 1    7.325687e-01
2    2.579158e-01
3    6.059978e-03
4    3.454862e-03
5    6.292246e-07
Name: payment_type, dtype: float64
```

```
In [33]: df=df[df['payment_type']<3]
df=df[(df['passenger_count'] >0) & (df['passenger_count']<6)]
```

```
In [ ]: df.shape
```

```
In [35]: df['payment_type'].replace([1,2],['Card','Cash'],inplace=True)
```

Descriptive statistics (describe helps to track outliers)

```
In [ ]: df.describe()
```

```
In [36]: df=df[df['passenger_count'] >0]
df=df[df['trip_distance'] >0]
df=df[df['fare_amount'] > 0]
```


Making countplot to identify general trends in the data

```
In [77]: ax = sns.countplot(x='payment_type',hue='passenger_count',data = df) # hue is used for kis basis pr chaie

for bars in ax.containers:
    ax.bar_label(bars)
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-77-e37c1b6ce327> in <module>
      2
      3 for bars in ax.containers:
----> 4     ax.bar_label(bars)

AttributeError: 'AxesSubplot' object has no attribute 'bar_label'
```




```
In [ ]: df['month']=df['reservation_date'].dt.month
```

```
plt.figure(figsize=(8,4))
ax1=sns.countplot(x='hotel',hue='is_canceled',data=df,palette='Blues')
legend_labels_ = ax1.get_legend_handles_labels()
ax1.legend(bbox_to_anchor(1,1))
plt.title('Reservation hotels in different city',size=20)
plt.xlabel=('No of hotels')
plt.ylabel=(['Not Canceled','Canceled'])
plt.show()
```


```
In [39]: # histogram and box plot helps in identifying outliers in this case
plt.hist(df['fare_amount'])
```

```
Out[39]: (array([2.2365779e+07, 1.8000000e+01, 1.0000000e+00, 0.0000000e+00,
0.0000000e+00, 0.0000000e+00, 1.0000000e+00, 0.0000000e+00,
0.0000000e+00, 1.0000000e+00]),
array([1.000000e-02, 6.964090e+02, 1.392808e+03, 2.089207e+03,
2.785606e+03, 3.482005e+03, 4.178404e+03, 4.874803e+03,
5.571202e+03, 6.267601e+03, 6.964000e+03]),
<BarContainer object of 10 artists>)
```



```
In [40]: plt.boxplot(df['fare_amount'])
```

```
Out[40]: {'whiskers': [<matplotlib.lines.Line2D at 0x26c93346520>,
<matplotlib.lines.Line2D at 0x26c93346880>],
'caps': [<matplotlib.lines.Line2D at 0x26c93346be0>,
<matplotlib.lines.Line2D at 0x26c93346f40>],
'boxes': [<matplotlib.lines.Line2D at 0x26c933461c0>],
'medians': [<matplotlib.lines.Line2D at 0x26c933572e0>],
'fliers': [<matplotlib.lines.Line2D at 0x26c93357640>],
'means': []}
```



Removing outlier using Z score or Interquartile range. Here we will use interquartile range

```
In [42]: for col in ['fare_amount', 'trip_distance', 'passenger_count']:
q1 = df[col].quantile(0.25)
q3 = df[col].quantile(0.75)
IQR = q3-q1

lower_bound = q1 - 1.5 * IQR
upper_bound = q3 + 1.5 * IQR

df=df[(df[col] >= lower_bound) & (df[col]<= upper_bound) ]
```

```
In [43]: df
```

```
Out[43]:
```


	passenger_count	payment_type	fare_amount	trip_distance	tpep_pickup_datetime	tpep_dropoff_datetime
0	1	Card	6.0	1.20	2020-01-01 00:28:15	2020-01-01 00:33:03
1	1	Card	7.0	1.20	2020-01-01 00:35:39	2020-01-01 00:43:04
2	1	Card	6.0	0.60	2020-01-01 00:47:41	2020-01-01 00:53:52
3	1	Card	5.5	0.80	2020-01-01 00:55:23	2020-01-01 01:00:14
5	1	Cash	2.5	0.03	2020-01-01 00:09:44	2020-01-01 00:10:37
...
24549230	1	Cash	6.0	1.20	2020-12-31 23:08:58	2020-12-31 23:14:28
24549233	1	Card	7.0	1.83	2020-12-31 22:55:17	2020-12-31 23:01:28
24549235	1	Card	9.0	2.18	2020-12-31 22:57:20	2020-12-31 23:05:33
24549236	1	Card	9.5	2.52	2020-12-31 23:40:35	2020-12-31 23:48:43
24549237	1	Card	4.5	0.59	2020-12-31 23:54:57	2020-12-31 23:57:39

14719854 rows × 6 columns

Fare amount and Trip distance analysis by payment mode


```
In [54]: plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.title('Distirbution of Fair Amount') # to give title
plt.hist(df[df['payment_type'] == 'Cash']['fare_amount'],histtype='barstacked',bins=20,edgecolor = 'k', color="blue",label='Cash')
plt.hist(df[df['payment_type'] == 'Card']['fare_amount'],histtype='barstacked',bins=20,edgecolor = 'k', color = "green",label='Card')
plt.legend() # plt.Legend shows the Label of the hist here cash and card
```

Out[54]: <matplotlib.legend.Legend at 0x26e141d3640>




```
In [55]: plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.title('Distirbution of Fair Amount') # to give title
plt.hist(df[df['payment_type'] == 'Cash']['fare_amount'],histtype='barstacked',bins=20,edgecolor = 'k', color="blue",label='Cash')
plt.hist(df[df['payment_type'] == 'Card']['fare_amount'],histtype='barstacked',bins=20,edgecolor = 'k', color = "green",label='Card')
plt.legend() # plt.Legend shows the Label of the hist here cash and card
```

Out[55]: <matplotlib.legend.Legend at 0x26e141eda90>




```
In [51]: plt.subplot(1,2,1)
plt.title('Distribution of Trip distance') # to give title
plt.hist(df[df['payment_type'] == 'Cash']['fare_amount'],histtype='barstacked',bins=20,edgecolor = 'k',color = 'blue',label='Cash')
plt.hist(df[df['payment_type'] == 'Card']['fare_amount'],histtype='barstacked',bins=20,edgecolor = 'k',color = 'green',label='Card')
plt.legend() # plt.Legend shows the Label of the hist here cash and card
plt.show()
```




```
In [57]: df.groupby('payment_type').agg({'fare_amount':['mean','std'],'trip_distance':['mean','std']})
```

Out[57]:

	fare_amount		trip_distance	
	mean	std	mean	std
payment_type				
Card	8.861820	3.575196	1.693476	0.966234
Cash	8.341731	3.628560	1.561408	0.991631

```
In [60]: plt.title('Preference of Payment Type')
plt.pie(df['payment_type'].value_counts(normalize=True), labels = df['payment_type'].value_counts().index,
        startangle=90,shadow=True,autopct='%1.1f%%',colors=['green','red'])
plt.show()
```



```
In [65]: passenger_count = df.groupby(['payment_type','passenger_count'])['passenger_count'].count()
passenger_count.rename(columns = {'passenger_count':'count'},inplace=True)
passenger_count.reset_index(inplace = True)
```


```
In [66]: passenger_count['perc'] = (passenger_count['count']/passenger_count['count'].sum()) *100
```

```
In [67]: passenger_count
```

Out[67]:

	payment_type	passenger_count	count	perc
0	Card	1	10985322	74.629286
1	Cash	1	3734532	25.370714

```
In [74]: sm.qqplot(df['fare_amount'],line='45')
plt.show()
```



```
In [69]: card_sample = df[df['payment_type']=='Card']['fare_amount']
cash_sample = df[df['payment_type']=='Cash']['fare_amount']
```

```
In [75]: t_stats,p_value = st.ttest_ind(a=card_sample, b=cash_sample, equal_var=False)
print('T_statistic',t_stats,'P-value',p_value)
```

T_statistic 240.17658261095323 P-value 0.0

```
In [ ]: cancelled_data = df[df['is_cancelled'] ==1]
top_10_country = cancelled_data['Country'].value_counts()[:10] #:10 sort data in descending count order by cancelled
plt.figure(figsize=(8,8))
plt.title('Top 10 countries with reservation cancelled')
plt.pie(top_10_country, autopct='.2f%', labels=top_10_country.index) #index returns the value based on top_10_country list
plt.show()
```

```
In [78]: ### Z test
age_mean = np.mean(age)
print(age_mean)
```

```
In [ ]: from statsmodels.stats import weightstats as stests

ztest,pvalue = stests.ztest(data['age'],value=30)
print(float(pvalue))
```

```
In [ ]: ## Anova
from scipy.stats import f_oneway
# Performance when each of the engine
# oil is applied
performance1 = [89, 89, 88, 78, 79]
performance2 = [93, 92, 94, 89, 88]
performance3 = [89, 88, 89, 93, 90]
performance4 = [81, 78, 81, 92, 82]

# Conduct the one-way ANOVA
f_oneway(performance1, performance2, performance3, performance4)
```

```
In [ ]: ### Chi square
from scipy.stats import chi2_contingency

# defining the table
data = [[207, 282, 241], [234, 242, 232]]
stat, p, dof, expected = chi2_contingency(data)

# interpret p-value
alpha = 0.05
print("p value is " + str(p))
if p <= alpha:
    print('Dependent (reject H0)')
else:
    print('Independent (H0 holds true)')
```